SnapClaim — print checklist

Day 2 (Oct 28 IST) — tick by hand. Freeze Nov 2. Submit Nov 4 eve IST. Worker: https://snapclaim-backend.parithosh.workers.dev

---

0. Hackathon frame

- [ ] Team of 3 locked: Parithosh, Tariq, Constantine
- [ ] Prize split in writing before submission (BLOCKS submission)
- [ ] Trial terms noted for Constantine
- [ ] Freeze Nov 2 — fixes only after
- [ ] Submit Nov 4 evening IST (Nov 4 9pm ET = Nov 5 7:30am IST)
- [ ] Today: Day 2 — Oct 28 IST / Oct 27 UTC
- [ ] Secrets stay local in .env (gitignored, never committed)
- [ ] Push only on explicit 'push' / 'ship' order
- [ ] No UI builds until freeze lifted (ACTIVE)

1. Deterministic core — DONE

- [ ] measureSession: masks in, findings + receipts out (src/engine/measure.ts:26)
- [ ] Math: areaRatio, calibratedLengthMm, mmPerPx, ratioToPercent
- [ ] Severity thresholds + classifySeverity + validateThresholds
- [ ] ReceiptLedger: every value has formula + inputs + confidence
- [ ] Evidence gate: uncited numbers fail with details[]
- [ ] SpendGuard: 429 at $0 remaining
- [ ] Mock adapters: backend invents no masks
- [ ] SVG annexure renderer, escaped, no PDF
- [ ] 40/40 tests green, tsc clean, $0 recorded

2. Backend router — DONE

- [ ] GET /v1/health -> {ok, version, time, mode}
- [ ] POST /v1/measure -> deterministic only, 50-mask cap
- [ ] POST /v1/report -> measure + findings + gate, 422 on fail
- [ ] POST /v1/segment -> mock ok; live 501, never fakes a mask
- [ ] POST /v1/findings -> gated drafts, usage ledgered in=/out=
- [ ] POST /v1/skeptic -> independent re-check, gated, 503 without key
- [ ] POST /v1/render -> gated report -> SVG {svg, gate, findings}
- [ ] GET /v1/spend -> {spentUsd, remainingUsd, limitUsd, entries}
- [ ] Enveloped replies + verbose errors + open CORS + quiet logs
- [ ] Limits: 50 masks/call, 1MB body, extra keys ignored
- [ ] Stateless replay; D1 stub throws until Day 3+
- [ ] Smoke green: live mode, gate 422 on tamper, $0
- [ ] D1 binding: wire or cut for submission

3. Worker deploy — DONE (mock, $0)

- [ ] wrangler.toml: snapclaim-backend, main worker.ts
- [ ] worker.ts imports handleRequest + corsHeaders, OPTIONS 204
- [ ] Deploy: npx wrangler deploy (backend dir)
- [ ] URL: https://snapclaim-backend.parithosh.workers.dev
- [ ] Verified: health mock; report 2 findings gate-ok absolute
- [ ] Submission demo link points at Worker URL
- [ ] Live keys on Worker: NOT done (mock default intentional)
- [ ] Rollback: wrangler rollback or re-deploy prior commit

4. Meta Spark findings — LIVE, pricing open

- [ ] Base https://api.meta.ai/v1, model muse-spark-1.3
- [ ] Path /v1/responses (chat path 404s — verified)
- [ ] Wall-crack smoke: 200, gate pass, 8 atomic findings
- [ ] Usage in=381 out=2658 logged at $0 (pricing unconfirmed)
- [ ] Confirm Meta per-token pricing before bulk runs
- [ ] Re-run on real photos, 1 call per photo, log each
- [ ] Keep $10 cap until submission week

5. SAM 3.1 masks — BLOCKED on shape (honest 501)

- [ ] /chat/completions -> 404 'Use /v1/responses' (verified)
- [ ] /images/segmentations, /segmentation, /embeddings -> 404
- [ ] /v1/responses -> 200 completed, images_processed:1, empty text x2
- [ ] /v1/segment stays 501 live; mock needs body.fixture
- [ ] Resolve shape via docs / verified tool def (no blind guessing)
- [ ] Wire vision provider, keep 501 fallback
- [ ] Add live-shape test only after shape known
- [ ] Re-test on real photos with coin/ruler
- [ ] Acceptance: real masks -> measure -> gate pass, zero synthetic

6. Baseten skeptic — CODE DONE, access blocked (403)

- [ ] Correct id meta-models/muse-spark-open (-1026 does not exist)
- [ ] Pricing verified: in $0.00000125 / out $0.00000425 per token
- [ ] SPARK_OPEN_PRICE in backend/providers.ts, charged via estimateCostUsd
- [ ] Independent prompt: re-derive only, mark DISPUTED with reason
- [ ] 503 without BASETEN_API_KEY; 429 at $0; mock-safe
- [ ] Tests: pricing math 0.003375, 503 no-key, mock gated drafts
- [ ] YOU: submit Baseten form / workspace grant (inference 403 now)
- [ ] Then 1 live /v1/skeptic smoke, log usage + spend
- [ ] Re-run on real-photo report; AGREED vs DISPUTED in narrative

7. Voice transcribe — PARKED (text first)

- [ ] Model listed: muse-voice-transcribe-1.0
- [ ] GET/POST /v1/audio/transcriptions -> 404 (verified)

- [ ] /responses + /chat/completions -> model_not_found (verified)
- [ ] Text input is primary; no voice UI, no transcribe call
- [ ] Resolve audio shape via docs/sample (no guessing)
- [ ] Wire /v1/transcribe mock-safe + test (after UI freeze lifted)

8. PWA — FROZEN (do not touch)

- [ ] Health, fixture masks JSON, transcript box
- [ ] Report + gate badges, schematic bars, raw JSON
- [ ] Skeptic AGREED/DISPUTED + citations, inline 503 hint
- [ ] Annexure SVG render + download
- [ ] Default backend = Worker URL (localhost typeable)
- [ ] NO voice button, styling, or new sections until lifted

9. Eval bench — LIVE (Constantine)

- [ ] eval/bench.ts + npm run eval (2 pass wall-crack baseline)
- [ ] fixtures/expected/wall-crack.json locked (5.17% 85.9mm sig; 1.29% null mod)
- [ ] Tolerance 0.005% / 0.051mm; missing expected = SKIP
- [ ] eval/README.md 3-step plug-in guide
- [ ] Constantine dataset 1: masks + expected (async PR)
- [ ] Gap table: model guess vs receipt value per mask
- [ ] Threshold sensitivity note (cosmetic/moderate/significant)
- [ ] Percent-only vs absolute mode cases

10. Real photos — MISSING (highest leverage)

- [ ] 3-5 spots: hairline crack, wide crack, stain/damp, peeling, 1 clean control
- [ ] Coin (Rs10, 27mm) or ruler taped flat beside damage, same plane
- [ ] 2 angles per spot: straight-on (measure) + 45 deg (context)
- [ ] Log per photo: time, indoor/outdoor, flash, phone model
- [ ] Distance 30-50cm, fill frame, no blur, no filter
- [ ] Do NOT vendor nipponpaint binary (URL-only control)
- [ ] Re-run measure -> findings -> skeptic per photo, log all

11. Spend ledger — $0 to date

- [ ] GET /v1/spend live; cap $10
- [ ] Spark smoke in=381 out=2658 at $0 (unconfirmed)
- [ ] Baseten pricing verified but unused
- [ ] Probes (404/403/empty) confirmed free
- [ ] Confirm Meta pricing before bulk runs
- [ ] Log every live call with in/out + usd

12. Team board

- [ ] You: Baseten form, photos, pricing, prize-split, submit call
- [ ] Tariq: next generalist slice (intake runs, copy draft, demo script)

- [ ] Constantine: eval dataset PR, gap table, async only
- [ ] GalaxyGamingFR invite accepted?
- [ ] Prize split signed before Nov 4

13. Marketing — PARKED

- [ ] Halted at Phase 1; no plan.md created
- [ ] Resume only on your order
- [ ] Gates: confirm between phases; public acts need explicit yes
- [ ] Never spend money; English copy

14. Risks

- [ ] SAM never resolves -> submit mock masks + real math + live text (say so)
- [ ] Baseten never granted -> skeptic stays mock-gated (say so)
- [ ] Transcribe never resolves -> text-only (already the plan)
- [ ] No real photos -> fixture-only demo (weakest; photos = top priority)
- [ ] Spend surprise -> caps on, log every call, no bulk runs

15. Definition of done

- [ ] Worker URL live, mode stated honestly
- [ ] 1+ real photo end-to-end: masks -> receipts -> findings -> gate -> SVG
- [ ] Gate demo: tampered draft -> 422 with details[]
- [ ] Spend page shows ledger; total in submission
- [ ] Eval output pasted (pass/skip/fail + gap table)
- [ ] README states live vs mock vs blocked (no overclaim)
- [ ] Demo script under 2 min: photo + 1 line -> report -> skeptic -> annexure
- [ ] Prize split signed, repo pushed, link filed Nov 4 eve IST

16. Commands

- [ ] npm test (40 green) / tsc --noEmit / npm run eval
- [ ] Mock: npm run backend:start -> :8787 / npm run backend:smoke
- [ ] Live: npm run backend:live (needs .env)
- [ ] Worker: npx wrangler deploy (backend dir)
- [ ] Git: status first; push only on 'push'/'ship'

Signed prize split: _____ Date: _____ Photos taken: [ ] Baseten form sent: [ ]